

数据库课程实践期末考试

院、系	计算机学院	年级专业	2020 软件工程	姓名	高歌	学号	2030416018
课程名称	数据库课程实践					成绩	

一. 补全 simpleDB 中的 TupleDesc.java 和 Tuple.java。

注：为了清晰起见，所有的“// some code goes here”都未删去，可根据这段注释定位代码。

(1) TupleDesc 的实现

首先补全 TupleDesc 相关代码。根据代码中给出的注释，可以了解 TupleDesc (TupleDescription) 负责描述某个一表中元组的各字段名以及类型信息。

考虑到 TupleDesc 需要存储 TDIItem，首先定义 items 属性：

```
private final TDIItem[] items;
```

然后实现构造器：

```
public TupleDesc(Type[] typeAr, String[] fieldAr) {
    // some code goes here
    items = new TDIItem[typeAr.length];
    for (int i = 0; i < typeAr.length; i++) {
        items[i] = new TDIItem(typeAr[i], fieldAr[i]);
    }
}
```

```
public TupleDesc(Type[] typeAr) {
    // some code goes here
    this(typeAr, new String[typeAr.length]);
}
```

然后实现几个基本方法：

```
public int numFields() {
    // some code goes here
    return items.length;
}

public String getFieldname(int i) throws NoSuchElementException {
    // some code goes here
    if (i < 0 || i >= items.length) {
        throw new NoSuchElementException();
    }
    return items[i].fieldName;
}

public Type getFieldType(int i) throws NoSuchElementException {
    // some code goes here
    if (i < 0 || i >= items.length) {
        throw new NoSuchElementException();
    }
}
```

```

        return items[i].fieldType;
    }

    public int fieldNameToIndex(String name) throws NoSuchElementException {
        // some code goes here
        for (int i = 0; i < items.length; i++) {
            if (items[i].fieldName != null && items[i].fieldName.equals(name)) {
                return i;
            }
        }
        throw new NoSuchElementException();
    }

    public int getSize() {
        // some code goes here
        int size = 0;
        for (int i = 0; i < items.length; i++) {
            size += items[i].fieldType.getLen();
        }
        return size;
    }

    public String toString() {
        // some code goes here
        StringBuilder sb = new StringBuilder();
        for (int i = 0; i < items.length; i++) {
            sb.append(items[i].fieldName);
            sb.append("(");
            sb.append(items[i].fieldType.toString());
            sb.append(")");
            if (i != items.length - 1) {
                sb.append(", ");
            }
        }
        return sb.toString();
    }
}

```

上面这些代码实现起来都比较简单，也没什么太多好谈的。下面将几个较为特殊的方法拿出来单独提一下。

首先是 `equals`，这是个容易实现但不太容易实现好的方法。一个通用的做法是，先与 `null` 值判断，接着判断是否与自身指向同一引用，再然后判断是否与自身属于同一个类，再往后才进行属性判断。子这里，要判断两个 `TupleDesc` 实例是否相等，先判断字段数是否相等，然后判断字段类型是否也依次相等（根据注释，这里的判断相等不要求字段名也对应相等）：

```

public boolean equals(Object o) {
    // some code goes here
    if (o == null) {

```

```

        return false;
    }
    if (o == this) {
        return true;
    }
    if (!(o instanceof TupleDesc)) {
        return false;
    }
    TupleDesc td = (TupleDesc) o;
    if (td.numFields() != numFields()) {
        return false;
    }
    for (int i = 0; i < numFields(); i++) {
        if (!getFieldType(i).equals(td.getFieldType(i))) {
            return false;
        }
    }
    return true;
}

```

另一个略微复杂的是 merge 的实现，但本质上也非常简单，即简单地创建新的 typeAr 和 fieldAr 将两者合并，然后创建新的 TupleDesc.

```

public static TupleDesc merge(TupleDesc td1, TupleDesc td2) {
    // some code goes here
    Type[] typeAr = new Type[td1.numFields() + td2.numFields()];
    String[] fieldAr = new String[td1.numFields() + td2.numFields()];
    for (int i = 0; i < td1.numFields(); i++) {
        typeAr[i] = td1.getFieldType(i);
        fieldAr[i] = td1.getFieldName(i);
    }
    for (int i = 0; i < td2.numFields(); i++) {
        typeAr[td1.numFields() + i] = td2.getFieldType(i);
        fieldAr[td1.numFields() + i] = td2.getFieldName(i);
    }
    return new TupleDesc(typeAr, fieldAr);
}

```

然后是标准的迭代器实现，比较简单，也没什么可多谈的：

```

public Iterator<TDItem> iterator() {
    // some code goes here
    return new TupleDescIterator();
}

private class TupleDescIterator implements Iterator<TDItem> {

    private int curr = 0;
}

```

```

@Override
public boolean hasNext() {
    return curr < numFields();
}

@Override
public TItem next() {
    return items[curr++];
}

@Override
public void remove() {
    throw new UnsupportedOperationException();
}
}

```

至此，TupleDesc 实现完毕。

(2) Tuple 的实现

根据注释，Tuple 负责存储表中的元组，它的模式由一个 TupleDesc 实例确定，并且包含一个可选的 RecordId 实例。因此为 Tuple 类添加下面三个实例属性：

```

private TupleDesc td;
private RecordId rid;
private final Field[] fields;

```

然后是构造器：

```

public Tuple(TupleDesc td) {
    // some code goes here
    this.td = td;
    this.rid = null;
    this.fields = new Field[td.numFields()];
}

```

接着是几个简单的方法，没什么好多说的：

```

public TupleDesc getTupleDesc() {
    // some code goes here
    return td;
}

public RecordId getRecordId() {
    // some code goes here
    return rid;
}

public void setRecordId(RecordId rid) {
    // some code goes here
    this.rid = rid;
}

```

```

public void setField(int i, Field f) {
    // some code goes here
    fields[i] = f;
}

public Field getField(int i) {
    // some code goes here
    return fields[i];
}

public String toString() {
    // some code goes here
    StringBuilder sb = new StringBuilder();
    for (int i = 0; i < fields.length; i++) {
        sb.append(fields[i].toString());
        if (i != fields.length - 1) {
            sb.append("\t");
        }
    }
    sb.append("\n");
    return sb.toString();
}

public void resetTupleDesc(TupleDesc td) {
    // some code goes here
    this.td = td;
}

```

然后是迭代器的实现，同 TupleDesc 区别不大：

```

public Iterator<Field> fields() {
    // some code goes here
    return new FieldIterator();
}

private class FieldIterator implements Iterator<Field> {
    private int i = 0;

    @Override
    public boolean hasNext() {
        // some code goes here
        return i < fields.length;
    }

    @Override
    public Field next() {

```

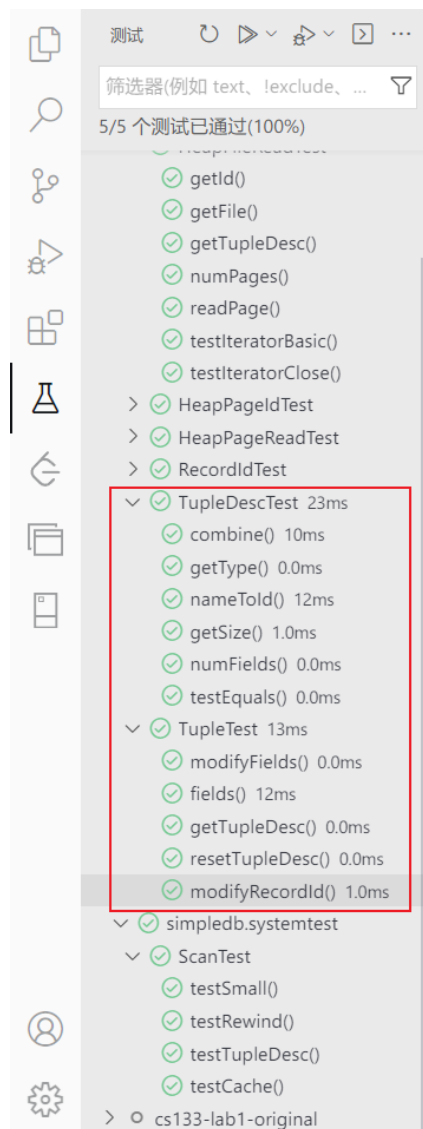
```

        // some code goes here
        return fields[i++];
    }

    @Override
    public void remove() {
        // some code goes here
        throw new UnsupportedOperationException();
    }
}

```

(3) 单元测试



*注：事实上写到这里时，所有代码都已经写完了，因此已经通过了全部测试

(4) 附：全部代码

```
// src/java/TupleDesc.java
```

```

package simpledb;

import java.io.Serializable;
import java.util.*;

/**
 * TupleDesc describes the schema of a tuple.
 */
public class TupleDesc implements Serializable {

    /**
     * A help class to facilitate organizing the information of each field
     */
    public static class TDIItem implements Serializable {

        private static final long serialVersionUID = 1L;

        /**
         * The type of the field
         */
        public final Type fieldType;

        /**
         * The name of the field
         */
        public final String fieldName;

        public TDIItem(Type t, String n) {
            this.fieldName = n;
            this.fieldType = t;
        }

        public String toString() {
            return fieldName + "(" + fieldType + ")";
        }
    }

    /**
     * @return
     *      An iterator which iterates over all the field TDIItems
     *      that are included in this TupleDesc
     */
    public Iterator<TDIItem> iterator() {
        // some code goes here
        return new TupleDescIterator();
    }
}

```

```

}

private class TupleDescIterator implements Iterator<TDItem> {

    private int curr = 0;

    public boolean hasNext() {
        return curr < numFields();
    }

    public TDItem next() {
        return items[curr++];
    }

    public void remove() {
        throw new UnsupportedOperationException();
    }
}

private static final long serialVersionUID = 1L;

private final TDItem[] items;

/**
 * Create a new TupleDesc with typeAr.length fields with fields of the
 * specified types, with associated named fields.
 *
 * @param typeAr
 *         array specifying the number of and types of fields in this
 *         TupleDesc. It must contain at least one entry.
 *
 * @param fieldAr
 *         array specifying the names of the fields. Note that names may
 *         be null.
 */
public TupleDesc(Type[] typeAr, String[] fieldAr) {
    // some code goes here
    items = new TDItem[typeAr.length];
    for (int i = 0; i < typeAr.length; i++) {
        items[i] = new TDItem(typeAr[i], fieldAr[i]);
    }
}

/**
 * Constructor. Create a new tuple desc with typeAr.length fields with
 * fields of the specified types, with anonymous (unnamed) fields.

```



```

*
* @param typeAr
*         array specifying the number of and types of fields in this
*         TupleDesc. It must contain at least one entry.
*/
public TupleDesc(Type[] typeAr) {
    // some code goes here
    this(typeAr, new String[typeAr.length]);
}

/**
* @return the number of fields in this TupleDesc
*/
public int numFields() {
    // some code goes here
    return items.length;
}

/**
* Gets the (possibly null) field name of the ith field of this TupleDesc.
*
* @param i
*         index of the field name to return. It must be a valid index.
* @return the name of the ith field
* @throws NoSuchElementException
*         if i is not a valid field reference.
*/
public String getFieldname(int i) throws NoSuchElementException {
    // some code goes here
    if (i < 0 || i >= items.length) {
        throw new NoSuchElementException();
    }
    return items[i].fieldName;
}

/**
* Gets the type of the ith field of this TupleDesc.
*
* @param i
*         The index of the field to get the type of. It must be a valid
*         index.
* @return the type of the ith field
* @throws NoSuchElementException
*         if i is not a valid field reference.
*/

```

```

public Type getFieldType(int i) throws NoSuchElementException {
    // some code goes here
    if (i < 0 || i >= items.length) {
        throw new NoSuchElementException();
    }
    return items[i].fieldType;
}

/**
 * Find the index of the field with a given name.
 *
 * @param name
 *         name of the field.
 * @return the index of the field that is first to have the given name.
 * @throws NoSuchElementException
 *         if no field with a matching name is found.
 */
public int fieldNameToIndex(String name) throws NoSuchElementException {
    // some code goes here
    for (int i = 0; i < items.length; i++) {
        if (items[i].fieldName != null && items[i].fieldName.equals(name)) {
            return i;
        }
    }
    throw new NoSuchElementException();
}

/**
 * @return The size (in bytes) of tuples corresponding to this TupleDesc.
 * Note that tuples from a given TupleDesc are of a fixed size.
 */
public int getSize() {
    // some code goes here
    int size = 0;
    for (int i = 0; i < items.length; i++) {
        size += items[i].fieldType.getLen();
    }
    return size;
}

/**
 * Merge two TupleDescs into one, with td1.numFields + td2.numFields fields,
 * with the first td1.numFields coming from td1 and the remaining from td2.
 *
 * @param td1

```

```

*           The TupleDesc with the first fields of the new TupleDesc
* @param td2
*           The TupleDesc with the last fields of the TupleDesc
* @return the new TupleDesc
*/
public static TupleDesc merge(TupleDesc td1, TupleDesc td2) {
    // some code goes here
    Type[] typeAr = new Type[td1.numFields() + td2.numFields()];
    String[] fieldAr = new String[td1.numFields() + td2.numFields()];
    for (int i = 0; i < td1.numFields(); i++) {
        typeAr[i] = td1.getFieldType(i);
        fieldAr[i] = td1.getFieldName(i);
    }
    for (int i = 0; i < td2.numFields(); i++) {
        typeAr[td1.numFields() + i] = td2.getFieldType(i);
        fieldAr[td1.numFields() + i] = td2.getFieldName(i);
    }
    return new TupleDesc(typeAr, fieldAr);
}

/**
 * Compares the specified object with this TupleDesc for equality. Two
 * TupleDescs are considered equal if they are the same size and if the n-th
 * type in this TupleDesc is equal to the n-th type in td.
 *
 * @param o
 *         the Object to be compared for equality with this TupleDesc.
 * @return true if the object is equal to this TupleDesc.
 */
public boolean equals(Object o) {
    // some code goes here
    if (o == null) {
        return false;
    }
    if (o == this) {
        return true;
    }
    if (!(o instanceof TupleDesc)) {
        return false;
    }
    TupleDesc td = (TupleDesc) o;
    if (td.numFields() != numFields()) {
        return false;
    }
    for (int i = 0; i < numFields(); i++) {

```

```

        if (!getFieldType(i).equals(td.getFieldType(i))) {
            return false;
        }
    }
    return true;
}

public int hashCode() {
    // If you want to use TupleDesc as keys for HashMap, implement this so
    // that equal objects have equals hashCode() results
    throw new UnsupportedOperationException("unimplemented");
}

/**
 * Returns a String describing this descriptor. It should be of the form
 * "fieldName[0](fieldType[0]), ..., fieldName[M](fieldType[M])", although
 * the exact format does not matter.
 *
 * @return String describing this descriptor.
 */
public String toString() {
    // some code goes here
    StringBuilder sb = new StringBuilder();
    for (int i = 0; i < items.length; i++) {
        sb.append(items[i].fieldName);
        sb.append("(");
        sb.append(items[i].fieldType.toString());
        sb.append(")");
        if (i != items.length - 1) {
            sb.append(", ");
        }
    }
    return sb.toString();
}
}

```

```
// src/java/Tuple.java
```

```

package simplifiedb;

import java.io.Serializable;
import java.util.Arrays;
import java.util.Iterator;

/**

```

```

* Tuple maintains information about the contents of a tuple. Tuples have a
* specified schema specified by a TupleDesc object and contain Field objects
* with the data for each field.
*/
public class Tuple implements Serializable {

    private static final long serialVersionUID = 1L;

    private TupleDesc td;
    private RecordId rid;
    private final Field[] fields;

    /**
     * Create a new tuple with the specified schema (type).
     *
     * @param td
     *         the schema of this tuple. It must be a valid TupleDesc
     *         instance with at least one field.
     */
    public Tuple(TupleDesc td) {
        // some code goes here
        this.td = td;
        this.rid = null;
        this.fields = new Field[td.numFields()];
    }

    /**
     * @return The TupleDesc representing the schema of this tuple.
     */
    public TupleDesc getTupleDesc() {
        // some code goes here
        return this.td;
    }

    /**
     * @return The RecordId representing the location of this tuple on disk. May
     *         be null.
     */
    public RecordId getRecordId() {
        // some code goes here
        return this.rid;
    }

    /**
     * Set the RecordId information for this tuple.

```

```

*
* @param rid
*         the new RecordId for this tuple.
*/
public void setRecordId(RecordId rid) {
    // some code goes here
    this.rid = rid;
}

/**
 * Change the value of the ith field of this tuple.
 *
 * @param i
 *         index of the field to change. It must be a valid index.
 * @param f
 *         new value for the field.
 */
public void setField(int i, Field f) {
    // some code goes here
    this.fields[i] = f;
}

/**
 * @return the value of the ith field, or null if it has not been set.
 *
 * @param i
 *         field index to return. Must be a valid index.
 */
public Field getField(int i) {
    // some code goes here
    return this.fields[i];
}

/**
 * Returns the contents of this Tuple as a string. Note that to pass the
 * system tests, the format needs to be as follows:
 *
 * column1\tcolumn2\tcolumn3\t...\tcolumnN\n
 *
 * where \t is any whitespace, except newline, and \n is a newline
 */
public String toString() {
    // some code goes here
    StringBuilder sb = new StringBuilder();
    for (int i = 0; i < this.fields.length; i++) {

```

```

        sb.append(this.fields[i].toString());
        if (i != this.fields.length - 1) {
            sb.append("\t");
        }
    }
    sb.append("\n");
    return sb.toString();
}

/**
 * @return
 *      An iterator which iterates over all the fields of this tuple
 */
public Iterator<Field> fields() {
    // some code goes here
    return new FieldIterator();
}

private class FieldIterator implements Iterator<Field> {
    private int i = 0;

    @Override
    public boolean hasNext() {
        // some code goes here
        return i < fields.length;
    }

    @Override
    public Field next() {
        // some code goes here
        return fields[i++];
    }

    @Override
    public void remove() {
        // some code goes here
        throw new UnsupportedOperationException();
    }
}

/**
 * Reset the TupleDesc of this tuple
 * Does not need to worry about the fields inside the Tuple
 */
public void resetTupleDesc(TupleDesc td) {

```

```

        // some code goes here
        this.td = td;
    }
}

```

二. 补全 simpleDB 中的 Catalog.java。

(1) Catalog.java 的实现

根据注释，每个数据库仅有一个 Catalog 实例，用于跟踪数据库中所有模式（即表）。对表的各类操作（添加、删除、修改、获取、查找文件等）都通过 Catalog 来实现。

首先考虑到 Catalog 用于进行对表的操作，因此必然存在一个 Map 对象用于存储表文件；此外考虑到 TupleDesc 并不包含主键相关信息，因此 Catalog 还需额外存储主键信息，这需要另一个 Map 对象。为线程安全考虑，均使用 ConcurrentMap 实现：

```

private final Map<String, DbFile> dbfiles;
private final Map<String, String> primaryKeys;

public Catalog() {
    // some code goes here
    dbfiles = new ConcurrentHashMap<String, DbFile>();
    primaryKeys = new ConcurrentHashMap<String, String>();
}

```

然后是添加表的方法，这里涉及了三次重载：

```

public void addTable(DbFile file, String name, String pkeyField) {
    // some code goes here
    if (name == null) {
        throw new IllegalArgumentException("name is null");
    }
    if (file == null) {
        throw new IllegalArgumentException("file is null");
    }
    if (pkeyField == null) {
        throw new IllegalArgumentException("pkeyField is null");
    }
    dbfiles.put(name, file);
    primaryKeys.put(name, pkeyField);
}

public void addTable(DbFile file, String name) {
    addTable(file, name, "");
}

public void addTable(DbFile file) {
    addTable(file, (UUID.randomUUID()).toString());
}

```

然后是一些简单的方法。为了简洁起见，这里部分使用了 Java 8 的 Stream API，因此在 Java 8 以

下版本应当无法运行。

```
public int getTableId(String name) throws NoSuchElementException {
    // some code goes here
    if (name == null) {
        throw new NoSuchElementException("name is null");
    }
    if (!dbfiles.containsKey(name)) {
        throw new NoSuchElementException("table " + name + " does not exist");
    }
    return dbfiles.get(name).getId();
}

public TupleDesc getTupleDesc(int tableid) throws NoSuchElementException {
    // some code goes here
    if (dbfiles.values().stream().noneMatch(x -> x.getId() == tableid)) {
        throw new NoSuchElementException("table " + tableid + " does not exist");
    }
    return dbfiles.get(getTableName(tableid)).getTupleDesc();
}

public DbFile getDatabaseFile(int tableid) throws NoSuchElementException {
    // some code goes here
    if (dbfiles.values().stream().noneMatch(x -> x.getId() == tableid)) {
        throw new NoSuchElementException("table " + tableid + " does not exist");
    }
    return dbfiles.get(getTableName(tableid));
}

public String getPrimaryKey(int tableid) {
    // some code goes here
    if (dbfiles.values().stream().noneMatch(x -> x.getId() == tableid)) {
        throw new NoSuchElementException("table " + tableid + " does not exist");
    }
    return primaryKeys.get(getTableName(tableid));
}

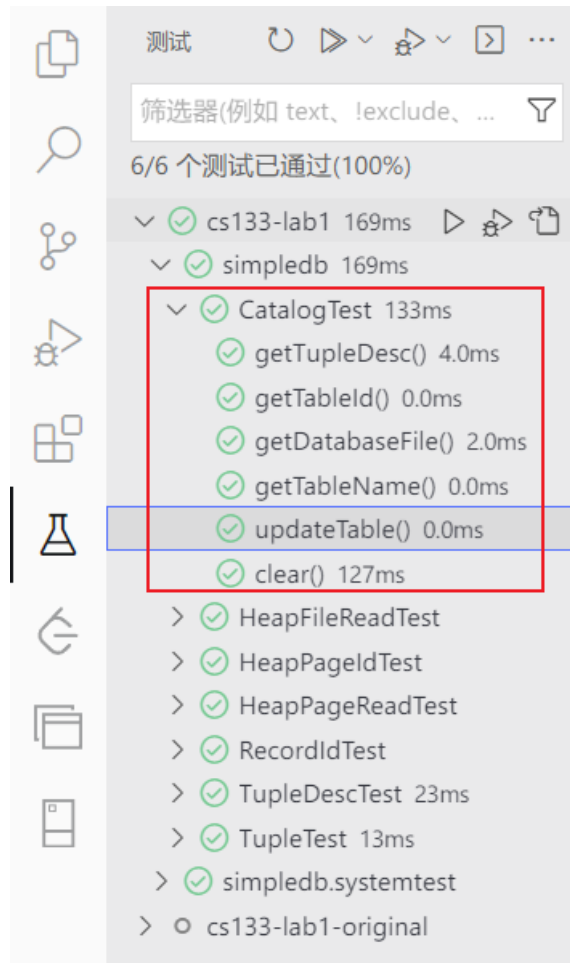
public String getTableName(int id) {
    // some code goes here
    return dbfiles.entrySet().stream()
        .filter(x -> x.getValue().getId() == id)
        .findFirst().get().getKey();
}

public void clear() {
    // some code goes here
    dbfiles.clear();
}

public Iterator<Integer> tableIdIterator() {
    // some code goes here
    return dbfiles.values().stream().map(x -> x.getId()).iterator();
}
```

}

(2) 单元测试



(3) 附：全部代码

```
// src/java/Catalog.java
```

```
package simpledb;
```

```
import java.io.BufferedReader;
```

```
import java.io.File;
```

```
import java.io.FileReader;
```

```
import java.io.IOException;
```

```
import java.util.*;
```

```
import java.util.concurrent.ConcurrentHashMap;
```

```
/**
```

```
 * The Catalog keeps track of all available tables in the database and their  
 * associated schemas.
```

```
 * For now, this is a stub catalog that must be populated with tables by a  
 * user program before it can be used -- eventually, this should be converted
```

```

* to a catalog that reads a catalog table from disk.
*
* @Threadsafe
*/
public class Catalog {

    private final Map<String, DbFile> dbfiles;
    private final Map<String, String> primaryKeys;

    /**
     * Constructor.
     * Creates a new, empty catalog.
     */
    public Catalog() {
        // some code goes here
        dbfiles = new ConcurrentHashMap<String, DbFile>();
        primaryKeys = new ConcurrentHashMap<String, String>();
    }

    /**
     * Add a new table to the catalog.
     * This table's contents are stored in the specified DbFile.
     *
     * @param file      the contents of the table to add; file.getId() is the
     *                  identifier of
     *                  this file/tupledesc param for the calls getTupleDesc and
     *                  getFile
     * @param name      the name of the table -- may be an empty string. May not be
     *                  null. If a name
     *                  conflict exists, use the last table to be added as the table
     *                  for a given name.
     * @param pkeyField the name of the primary key field
     */
    public void addTable(DbFile file, String name, String pkeyField) {
        // some code goes here
        if (name == null) {
            throw new IllegalArgumentException("name is null");
        }
        if (file == null) {
            throw new IllegalArgumentException("file is null");
        }
        if (pkeyField == null) {
            throw new IllegalArgumentException("pkeyField is null");
        }
        dbfiles.put(name, file);
    }
}

```

```

        primaryKeys.put(name, pkeyField);
    }

    public void addTable(DbFile file, String name) {
        addTable(file, name, "");
    }

    /**
     * Add a new table to the catalog.
     * This table has tuples formatted using the specified TupleDesc and its
     * contents are stored in the specified DbFile.
     *
     * @param file the contents of the table to add; file.getId() is the identifier
     *            of
     *            this file/tupledesc param for the calls getTupleDesc and getFile
     */
    public void addTable(DbFile file) {
        addTable(file, (UUID.randomUUID()).toString());
    }

    /**
     * Return the id of the table with a specified name,
     *
     * @throws NoSuchElementException if the table doesn't exist
     */
    public int getTableId(String name) throws NoSuchElementException {
        // some code goes here
        if (name == null) {
            throw new NoSuchElementException("name is null");
        }
        if (!dbfiles.containsKey(name)) {
            throw new NoSuchElementException("table " + name + " does not exist");
        }
        return dbfiles.get(name).getId();
    }

    /**
     * Returns the tuple descriptor (schema) of the specified table
     *
     * @param tableid The id of the table, as specified by the DbFile.getId()
     *                function passed to addTable
     * @throws NoSuchElementException if the table doesn't exist
     */
    public TupleDesc getTupleDesc(int tableid) throws NoSuchElementException {
        // some code goes here
    }

```

```

        if (dbfiles.values().stream().noneMatch(x -> x.getId() == tableid)) {
            throw new NoSuchElementException("table " + tableid + " does not exist");
        }
        return dbfiles.get(getTableName(tableid)).getTupleDesc();
    }

    /**
     * Returns the DbFile that can be used to read the contents of the
     * specified table.
     *
     * @param tableid The id of the table, as specified by the DbFile.getId()
     *                 function passed to addTable
     */
    public DbFile getDatabaseFile(int tableid) throws NoSuchElementException {
        // some code goes here
        if (dbfiles.values().stream().noneMatch(x -> x.getId() == tableid)) {
            throw new NoSuchElementException("table " + tableid + " does not exist");
        }
        return dbfiles.get(getTableName(tableid));
    }

    public String getPrimaryKey(int tableid) {
        // some code goes here
        if (dbfiles.values().stream().noneMatch(x -> x.getId() == tableid)) {
            throw new NoSuchElementException("table " + tableid + " does not exist");
        }
        return primaryKeys.get(getTableName(tableid));
    }

    public Iterator<Integer> tableIdIterator() {
        // some code goes here
        return dbfiles.values().stream().map(x -> x.getId()).iterator();
    }

    public String getTableName(int id) {
        // some code goes here
        return dbfiles.entrySet().stream()
            .filter(x -> x.getValue().getId() == id)
            .findFirst().get().getKey();
    }

    /** Delete all tables from the catalog */
    public void clear() {
        // some code goes here
        dbfiles.clear();
    }

```

```

}

/**
 * Reads the schema from a file and creates the appropriate tables in the
 * database.
 *
 * @param catalogFile
 */
public void loadSchema(String catalogFile) {
    String line = "";
    String baseFolder = new File(new File(catalogFile).getAbsolutePath())
        .getParent();
    try {
        BufferedReader br = new BufferedReader(
            new FileReader(new File(catalogFile)));

        while ((line = br.readLine()) != null) {
            // assume line is of the format name (field type, field type, ...)
            String name = line.substring(0, line.indexOf("(")).trim();
            // System.out.println("TABLE NAME: " + name);
            String fields = line.substring(line.indexOf("(") + 1,
line.indexOf(")").trim());
            String[] els = fields.split(",");
            ArrayList<String> names = new ArrayList<String>();
            ArrayList<Type> types = new ArrayList<Type>();
            String primaryKey = "";
            for (String e : els) {
                String[] els2 = e.trim().split(" ");
                names.add(els2[0].trim());
                if (els2[1].trim().toLowerCase().equals("int"))
                    types.add(Type.INT_TYPE);
                else if (els2[1].trim().toLowerCase().equals("string"))
                    types.add(Type.STRING_TYPE);
                else {
                    System.out.println("Unknown type " + els2[1]);
                    System.exit(0);
                }
                if (els2.length == 3) {
                    if (els2[2].trim().equals("pk"))
                        primaryKey = els2[0].trim();
                    else {
                        System.out.println("Unknown annotation " + els2[2]);
                        System.exit(0);
                    }
                }
            }
        }
    }
}

```

```

    }
    Type[] typeAr = types.toArray(new Type[0]);
    String[] namesAr = names.toArray(new String[0]);
    TupleDesc t = new TupleDesc(typeAr, namesAr);
    HeapFile tabHf = new HeapFile(new File(baseFolder + "/" + name +
".dat"), t);

    addTable(tabHf, name, primaryKey);
    System.out.println("Added table : " + name + " with schema " + t);
}
} catch (IOException e) {
    e.printStackTrace();
    System.exit(0);
} catch (IndexOutOfBoundsException e) {
    System.out.println("Invalid catalog entry : " + line);
    System.exit(0);
}
}
}
}

```

三. 补全 simpleDB 中的 BufferPool.java。

(1) BufferPool.java 的实现

根据注释，BufferPool 用于缓冲，访问数据时首先通过缓冲池获得，若不存在则加入缓冲池。当缓冲池已满时，丢弃第一个页面，若该页面已被更改，则先保存后再丢弃。

因此，需要包含一个 Map 对象用于存储所有 Page，还需要 numPages 变量用于保存 Page 的最大数量。此外，为了在缓冲池满时丢弃页面，还需要一个队列用于存储 pageId 加入的顺序，这里使用数组+一个用于存储数组下标的变量模拟这个队列：

```

private final ConcurrentHashMap<PageId, Page> buffer;
private final int numPages;
private final PageId[] pageIds;
private int pageIdIndex;

public BufferPool(int numPages) {
    // some code goes here
    buffer = new ConcurrentHashMap<PageId, Page>(numPages);
    this.numPages = numPages;
    pageIds = new PageId[numPages];
    pageIdIndex = 0;
}

```

然后根据题目，补全 getPage()方法即可：

```

public Page getPage(TransactionId tid, PageId pid, Permissions perm)
    throws TransactionAbortedException, DbException, IOException {
    // some code goes here
    if (tid == null) {
        throw new DbException("Transaction is null");
    }
}

```

```

    }
    if (pid == null) {
        throw new DbException("PageId is null");
    }
    if (perm == null) {
        throw new DbException("Permissions is null");
    }
    if (perm != Permissions.READ_ONLY && perm != Permissions.READ_WRITE) {
        throw new DbException("Permissions is not READ_ONLY or READ_WRITE");
    }
    if (buffer.containsKey(pid)) {
        return buffer.get(pid);
    } else {
        Page page = Database.getCatalog().getDatabaseFile(pid.getTableId())
            .readPage(pid);
        if (page == null) {
            throw new DbException("Page is null");
        }
        if (buffer.size() == numPages) {
            if (buffer.get(pageIds[0]).isDirty() != null) {
                flushPage(pageIds[0]);
            }
            buffer.remove(pageIds[0]);
            for (int i = 0; i < numPages - 1; i++) {
                pageIds[i] = pageIds[i + 1];
            }
            pageIds[numPages - 1] = pid;
        } else {
            pageIds[pageIdIndex] = pid;
            pageIdIndex++;
        }
        buffer.put(pid, page);
        return page;
    }
}

```

(2) 单元测试

本题不包含对应的单元测试。

(3) 附：全部代码

```
// src/java/BufferPool.java
```

```
package simplifiedb;
```

```
import java.io.*;
```



```

import java.util.concurrent.ConcurrentHashMap;

/**
 * BufferPool manages the reading and writing of pages into memory from
 * disk. Access methods call into it to retrieve pages, and it fetches
 * pages from the appropriate location.
 * <p>
 * The BufferPool is also responsible for locking; when a transaction fetches
 * a page, BufferPool checks that the transaction has the appropriate
 * locks to read/write the page.
 *
 * @Threadsafe, all fields are final
 */
public class BufferPool {
    /** Bytes per page, including header. */
    public static final int PAGE_SIZE = 4096;

    private static int pageSize = PAGE_SIZE;

    /**
     * Default number of pages passed to the constructor. This is used by
     * other classes. BufferPool should use the numPages argument to the
     * constructor instead.
     */
    public static final int DEFAULT_PAGES = 50;

    private final ConcurrentHashMap<PageId, Page> buffer;
    private final int numPages;

    /**
     * TODO for Lab 4: create your private Lock Manager class.
     * Be sure to instantiate it in the constructor.
     */

    /**
     * Creates a BufferPool that caches up to numPages pages.
     *
     * @param numPages maximum number of pages in this buffer pool.
     */
    public BufferPool(int numPages) {
        // some code goes here
        buffer = new ConcurrentHashMap<PageId, Page>(numPages);
        this.numPages = numPages;
    }
}

```

```

public static int getPageSize() {
    return pageSize;
}

/**
 * Helper: this should be used for testing only!!!
 */
public static void setPageSize(int pageSize) {
    BufferPool.pageSize = pageSize;
}

/**
 * Retrieve the specified page with the associated permissions.
 * Will acquire a lock and may block if that lock is held by another
 * transaction.
 * <p>
 * The retrieved page should be looked up in the buffer pool. If it
 * is present, it should be returned. If it is not present, it should
 * be added to the buffer pool and returned. If there is insufficient
 * space in the buffer pool, a page should be evicted and the new page
 * should be added in its place.
 *
 * @param tid the ID of the transaction requesting the page
 * @param pid the ID of the requested page
 * @param perm the requested permissions on the page
 */
public Page getPage(TransactionId tid, PageId pid, Permissions perm)
    throws TransactionAbortedException, DbException {
    // some code goes here
    if (tid == null) {
        throw new DbException("Transaction is null");
    }
    if (pid == null) {
        throw new DbException("PageId is null");
    }
    if (perm == null) {
        throw new DbException("Permissions is null");
    }
    if (perm != Permissions.READ_ONLY && perm != Permissions.READ_WRITE) {
        throw new DbException("Permissions is not READ_ONLY or READ_WRITE");
    }
    if (buffer.containsKey(pid)) {
        return buffer.get(pid);
    } else {

```

```

        Page page = Database.getCatalog().getDatabaseFile(pid.getTableId())
            .readPage(pid);
        if (page == null) {
            throw new DbException("Page is null");
        }
        if (buffer.size() == numPages) {
            throw new DbException("BufferPool is full");
        }
        buffer.put(pid, page);
        return page;
    }
}

/**
 * Releases the lock on a page.
 * Calling this is very risky, and may result in wrong behavior. Think hard
 * about who needs to call this and why, and why they can run the risk of
 * calling it.
 *
 * @param tid the ID of the transaction requesting the unlock
 * @param pid the ID of the page to unlock
 */
public void releasePage(TransactionId tid, PageId pid) {
    // some code goes here
    // not necessary for Lab1/Lab2
}

/**
 * Release all locks associated with a given transaction.
 *
 * @param tid the ID of the transaction requesting the unlock
 */
public void transactionComplete(TransactionId tid) throws IOException {
    // some code goes here
    // not necessary for Lab1/Lab2
}

/** Return true if the specified transaction has a lock on the specified page
 */
public boolean holdsLock(TransactionId tid, PageId p) {
    // some code goes here
    // not necessary for Lab1/Lab2
    return false;
}
}

```

```

/**
 * Commit or abort a given transaction; release all locks associated to
 * the transaction.
 *
 * @param tid    the ID of the transaction requesting the unlock
 * @param commit a flag indicating whether we should commit or abort
 */
public void transactionComplete(TransactionId tid, boolean commit)
    throws IOException {
    // some code goes here
    // not necessary for lab1/lab2
}

/**
 * Add a tuple to the specified table on behalf of transaction tid. Will
 * acquire a write lock on the page the tuple is added to and any other
 * pages that are updated (Lock acquisition is not needed for Lab2).
 * May block if the lock(s) cannot be acquired.
 *
 * Marks any pages that were dirtied by the operation as dirty by calling
 * their markDirty bit, and updates cached versions of any pages that have
 * been dirtied so that future requests see up-to-date pages.
 *
 * @param tid    the transaction adding the tuple
 * @param tableId the table to add the tuple to
 * @param t      the tuple to add
 */
public void insertTuple(TransactionId tid, int tableId, Tuple t)
    throws DbException, IOException, TransactionAbortedException {
    // some code goes here
    // not necessary for lab1
}

/**
 * Remove the specified tuple from the buffer pool.
 * Will acquire a write lock on the page the tuple is removed from and any
 * other pages that are updated. May block if the lock(s) cannot be acquired.
 *
 * Marks any pages that were dirtied by the operation as dirty by calling
 * their markDirty bit, and updates cached versions of any pages that have
 * been dirtied so that future requests see up-to-date pages.
 *
 * @param tid the transaction deleting the tuple.
 * @param t   the tuple to delete
 */

```

```

public void deleteTuple(TransactionId tid, Tuple t)
    throws DbException, IOException, TransactionAbortedException {
    // some code goes here
    // not necessary for Lab1
}

/**
 * Flush all dirty pages to disk.
 * NB: Be careful using this routine -- it writes dirty data to disk so will
 * break simpledb if running in NO STEAL mode.
 */
public synchronized void flushAllPages() throws IOException {
    // some code goes here
    // not necessary for Lab1
}

/**
 * Remove the specific page id from the buffer pool.
 * Needed by the recovery manager to ensure that the
 * buffer pool doesn't keep a rolled back page in its
 * cache.
 */
public synchronized void discardPage(PageId pid) {
    // some code goes here
    // not necessary for Labs 1--4
}

/**
 * Flushes a certain page to disk
 *
 * @param pid an ID indicating the page to flush
 */
private synchronized void flushPage(PageId pid) throws IOException {
    // some code goes here
    // not necessary for Lab1
}

/**
 * Write all pages of the specified transaction to disk.
 */
public synchronized void flushPages(TransactionId tid) throws IOException {
    // some code goes here
    // not necessary for Labs 1--4
}

```

```

/**
 * Discards a page from the buffer pool.
 * Flushes the page to disk to ensure dirty pages are updated on disk.
 */
private synchronized void evictPage() throws DbException {
    // some code goes here
    // not necessary for Lab1
}
}

```

四. 补全 simpleDB 中的 HeapPageId.java、RecordId.java、HeapPage.java。

(1) HeapPageId.java 与 RecordId.java 的实现

两个 Id 类的实现都比较简单，大同小异，这里就放在一起了。

首先是构造器，根据构造器接受的参数可直接定义所需要的属性：

```

private final int tableid;
private final int pgNo;

public HeapPageId(int tableId, int pgNo) {
    // some code goes here
    this.tableid = tableId;
    this.pgNo = pgNo;
}

```

```

private final PageId pageid;
private final int tupleno;

public RecordId(PageId pid, int tupleno) {
    // some code goes here
    this.pageid = pid;
    this.tupleno = tupleno;
}

```

然后是一些简单的 get/set 方法，大同小异，也没什么可说的：

```

public int getTableId() {
    // some code goes here
    return tableid;
}

public int pageNumber() {
    // some code goes here
    return pgNo;
}

```

```

public int tupleno() {

```

```

        // some code goes here
        return tupleno;
    }

```

```

    public PageId getPageId() {
        // some code goes here
        return pageid;
    }

```

然后为了能够放入 Map 中，还需要实现 hashCode()方法。这里的实现比较简单，只是单纯地将两个属性组合：

```

    public int hashCode() {
        // some code goes here
        return tableid * 1000 + pgNo;
    }

```

```

    public int hashCode() {
        // some code goes here
        return this.pageid.hashCode() + this.tupleno;
    }

```

接着是 equals 方法，这个上面已经提过了，也不赘述：

```

    public boolean equals(Object o) {
        // some code goes here
        if (o == null) {
            return false;
        }
        if (o == this) {
            return true;
        }
        if (!(o instanceof HeapPageId)) {
            return false;
        }
        HeapPageId pid = (HeapPageId) o;
        if (this.tableid == pid.tableid && this.pgNo == pid.pgNo) {
            return true;
        }
        return false;
    }

```

```

    public boolean equals(Object o) {
        // some code goes here
        if (o == null) {
            return false;
        }
        if (o == this) {
            return true;
        }

```

```

    }
    if (!(o instanceof RecordId)) {
        return false;
    }
    RecordId rid = (RecordId) o;
    if (this.pageid.equals(rid.pageid) && this.tupleno == rid.tupleno) {
        return true;
    }
    return false;
}

```

(2) HeapPage.java 的实现

根据注释，表文件存储在堆上（HeapFile），而一个 HeapFile 包含若干页面（HeapPage），一个 HeapPage 包含若干 slot，每个 slot 可以容纳一个元组。在本题中仅需实现 HeapPage.java。

尽管 HeapPage 的实现比较复杂，但由于大多数代码已经补全，因此只需补全少量方法。

首先是一些简单的 get/set 方法：

```

private int getNumTuples() {
    // some code goes here
    return (int) Math.floor((BufferPool.getPageSize() * 8)
        / (td.getSize() * 8 + 1));
}

private int getHeaderSize() {
    // some code goes here
    return (int) Math.ceil(numSlots / 8.0);
}

public HeapPageId getId() {
    // some code goes here
    return pid;
}

public int getNumEmptySlots() {
    // some code goes here
    int count = 0;
    for (int i = 0; i < tuples.length; i++) {
        if (!isSlotUsed(i)) {
            count++;
        }
    }
    return count;
}

public boolean isSlotUsed(int i) {
    // some code goes here
    return (header[i / 8] & (1 << (i % 8))) != 0;
}

```

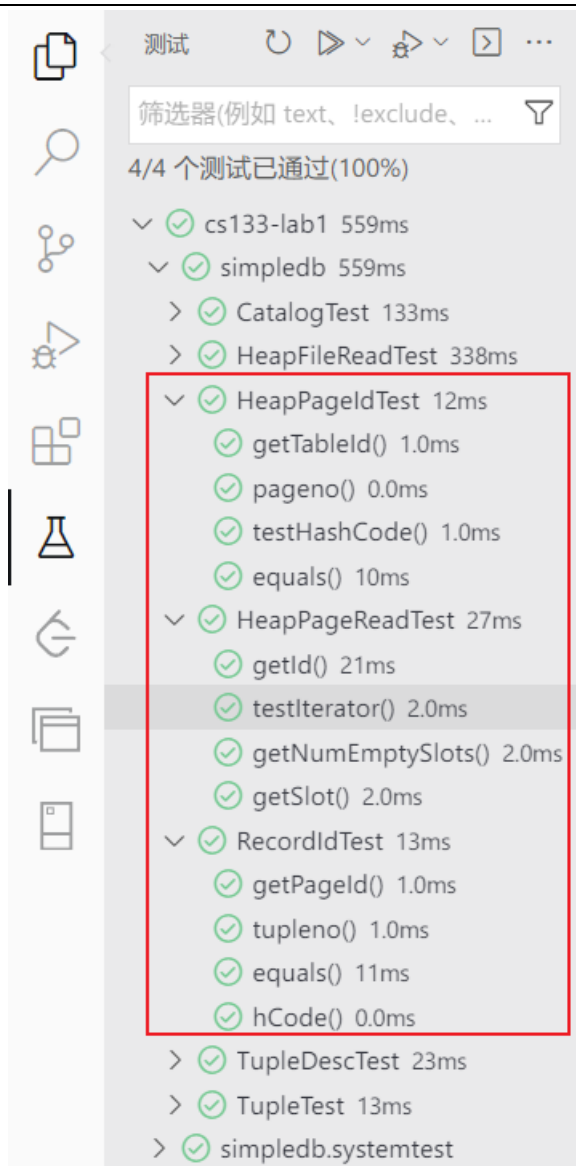
值得注意的是，尽管这些方法代码量很少，但确实需要一番思考。例如 getNumTuples 需要根据 Page 所占的大小除以 TupleDesc 所占大小得到；判断某个 slot 是否被使用需要在 header 数组中精准定

位索引指向的二进制位，由于 header 是一个 Byte 数组，因此每个单元都可以存储八个 slot 的状态，需要较为复杂的判断。

然后是迭代器，这个比较常规，就不多提了。需要注意的仅有判断 hasNext 时还需要判断一下 isSlotUsed.

```
public Iterator<Tuple> iterator() {  
    // some code goes here  
    return new HeapPageIterator();  
}  
  
private class HeapPageIterator implements Iterator<Tuple> {  
  
    private int curr = 0;  
  
    public boolean hasNext() {  
        return curr < tuples.length && isSlotUsed(curr);  
    }  
  
    public Tuple next() {  
        if (!hasNext()) {  
            throw new NoSuchElementException();  
        }  
        Tuple t = tuples[curr];  
        curr++;  
        return t;  
    }  
  
    public void remove() {  
        throw new UnsupportedOperationException();  
    }  
}
```

(3) 单元测试



(4) 附：全部代码

```
// src/java/HeapPageId.java
```

```
package simpledb;
```

```
/** Unique identifier for HeapPage objects. */
```

```
public class HeapPageId implements PageId {
```

```
    private final int tableid;
```

```
    private final int pgNo;
```

```
    /**
```

```
     * Constructor. Create a page id structure for a specific page of a  
     * specific table.
```

```
     *
```

```
     * @param tableId The table that is being referenced
```

```

    * @param pgNo    The page number in that table.
    */
    public HeapPageId(int tableId, int pgNo) {
        // some code goes here
        this.tableid = tableId;
        this.pgNo = pgNo;
    }

    /** @return the table associated with this PageId */
    public int getTableId() {
        // some code goes here
        return tableid;
    }

    /**
     * @return the page number in the table getTableId() associated with
     *         this PageId
     */
    public int pageNumber() {
        // some code goes here
        return pgNo;
    }

    /**
     * @return a hash code for this page, represented by the concatenation of
     *         the table number and the page number (needed if a PageId is used as
a
     *         key in a hash table in the BufferPool, for example.)
     * @see BufferPool
     */
    public int hashCode() {
        // some code goes here
        return tableid * 1000 + pgNo;
    }

    /**
     * Compares one PageId to another.
     *
     * @param o The object to compare against (must be a PageId)
     * @return true if the objects are equal (e.g., page numbers and table
     *         ids are the same)
     */
    public boolean equals(Object o) {
        // some code goes here
        if (o == null) {

```

```

        return false;
    }
    if (o == this) {
        return true;
    }
    if (!(o instanceof HeapPageId)) {
        return false;
    }
    HeapPageId pid = (HeapPageId) o;
    if (this.tableid == pid.tableid && this.pgNo == pid.pgNo) {
        return true;
    }
    return false;
}

/**
 * Return a representation of this object as an array of
 * integers, for writing to disk. Size of returned array must contain
 * number of integers that corresponds to number of args to one of the
 * constructors.
 */
public int[] serialize() {
    int data[] = new int[2];

    data[0] = getTableId();
    data[1] = pageNumber();

    return data;
}
}

```

// src/java/RecordId.java

```

package simpledb;

import java.io.Serializable;

/**
 * A RecordId is a reference to a specific tuple on a specific page of a
 * specific table.
 */
public class RecordId implements Serializable {

    private static final long serialVersionUID = 1L;
}

```

```

private final PageId pageid;
private final int tupleno;

/**
 * Creates a new RecordId referring to the specified PageId and tuple
 * number.
 *
 * @param pid
 *         the pageid of the page on which the tuple resides
 * @param tupleno
 *         the tuple number within the page.
 */
public RecordId(PageId pid, int tupleno) {
    // some code goes here
    this.pageid = pid;
    this.tupleno = tupleno;
}

/**
 * @return the tuple number this RecordId references.
 */
public int tupleno() {
    // some code goes here
    return tupleno;
}

/**
 * @return the page id this RecordId references.
 */
public PageId getPageId() {
    // some code goes here
    return pageid;
}

/**
 * Two RecordId objects are considered equal if they represent the same
 * tuple.
 *
 * @return True if this and o represent the same tuple
 */
@Override
public boolean equals(Object o) {
    // some code goes here
    if (o == null) {
        return false;
    }

```

```

    }
    if (o == this) {
        return true;
    }
    if (!(o instanceof RecordId)) {
        return false;
    }
    RecordId rid = (RecordId) o;
    if (this.pageid.equals(rid.pageid) && this.tupleno == rid.tupleno) {
        return true;
    }
    return false;
}

/**
 * You should implement the hashCode() so that two equal RecordId instances
 * (with respect to equals()) have the same hashCode().
 *
 * @return An int that is the same for equal RecordId objects.
 */
@Override
public int hashCode() {
    // some code goes here
    return this.pageid.hashCode() + this.tupleno;
}
}

```

```
// src/java/HeapPage.java
```

```

package simpledb;

import java.util.*;
import java.io.*;

/**
 * Each instance of HeapPage stores data for one page of HeapFiles and
 * implements the Page interface that is used by BufferPool.
 *
 * @see HeapFile
 * @see BufferPool
 */
public class HeapPage implements Page {

    final HeapPageId pid;

```

```

final TupleDesc td;
final byte header[];
final Tuple tuples[];
final int numSlots;

byte[] oldData;
private final Byte oldDataLock = new Byte((byte) 0);

/**
 * Create a HeapPage from a set of bytes of data read from disk.
 * The format of a HeapPage is a set of header bytes indicating
 * the slots of the page that are in use, some number of tuple slots.
 * Specifically, the number of tuples is equal to:
 * <p>
 *  $\text{floor}((\text{BufferPool.getPageSize()} * 8) / (\text{tuple size} * 8 + 1))$ 
 * <p>
 * where tuple size is the size of tuples in this
 * database table, which can be determined via {@link Catalog#getTupleDesc}.
 * The number of 8-bit header words is equal to:
 * <p>
 *  $\text{ceiling}(\text{no. tuple slots} / 8)$ 
 * <p>
 *
 * @see Database#getCatalog
 * @see Catalog#getTupleDesc
 * @see BufferPool#getPageSize()
 */
public HeapPage(HeapPageId id, byte[] data) throws IOException {
    this.pid = id;
    this.td = Database.getCatalog().getTupleDesc(id.getTableId());
    this.numSlots = getNumTuples();
    DataInputStream dis = new DataInputStream(new ByteArrayInputStream(data));

    // allocate and read the header slots of this page
    header = new byte[getHeaderSize()];
    for (int i = 0; i < header.length; i++)
        header[i] = dis.readByte();

    tuples = new Tuple[numSlots];
    try {
        // allocate and read the actual records of this page
        for (int i = 0; i < tuples.length; i++)
            tuples[i] = readNextTuple(dis, i);
    } catch (NoSuchElementException e) {
        e.printStackTrace();
    }
}

```

```

    }
    dis.close();

    setBeforeImage();
}

/**
 * Retrieve the number of tuples on this page.
 *
 * @return the number of tuples on this page
 */
private int getNumTuples() {
    // some code goes here
    return (int) Math.floor((BufferPool.getPageSize() * 8) / (td.getSize() * 8
+ 1));
}

/**
 * Computes the number of bytes in the header of a page in a HeapFile with each
 * tuple occupying tupleSize bytes
 *
 * @return the number of bytes in the header of a page in a HeapFile with each
 *         tuple occupying tupleSize bytes
 */
private int getHeaderSize() {
    // some code goes here
    return (int) Math.ceil(numSlots / 8.0);
}

/**
 * Return a view of this page before it was modified
 * -- used by recovery
 */
public HeapPage getBeforeImage() {
    try {
        byte[] oldDataRef = null;
        synchronized (oldDataLock) {
            oldDataRef = oldData;
        }
        return new HeapPage(pid, oldDataRef);
    } catch (IOException e) {
        e.printStackTrace();
        // should never happen -- we parsed it OK before!
        System.exit(1);
    }
}

```



```

        return null;
    }

    public void setBeforeImage() {
        synchronized (oldDataLock) {
            oldData = getPageData().clone();
        }
    }

    /**
     * @return the PageId associated with this page.
     */
    public HeapPageId getId() {
        // some code goes here
        return pid;
    }

    /**
     * Suck up tuples from the source file.
     */
    private Tuple readNextTuple(DataInputStream dis, int slotId) throws
    NoSuchElementException {
        // if associated bit is not set, read forward to the next tuple, and
        // return null.
        if (!isSlotUsed(slotId)) {
            for (int i = 0; i < td.getSize(); i++) {
                try {
                    dis.readByte();
                } catch (IOException e) {
                    throw new NoSuchElementException("error reading empty tuple");
                }
            }
            return null;
        }

        // read fields in the tuple
        Tuple t = new Tuple(td);
        RecordId rid = new RecordId(pid, slotId);
        t.setRecordId(rid);
        try {
            for (int j = 0; j < td.numFields(); j++) {
                Field f = td.getFieldType(j).parse(dis);
                t.setField(j, f);
            }
        } catch (java.text.ParseException e) {

```

```

        e.printStackTrace();
        throw new NoSuchElementException("parsing error!");
    }

    return t;
}

/**
 * Generates a byte array representing the contents of this page.
 * Used to serialize this page to disk.
 * <p>
 * The invariant here is that it should be possible to pass the byte
 * array generated by getPageData to the HeapPage constructor and
 * have it produce an identical HeapPage object.
 *
 * @see #HeapPage
 * @return A byte array correspond to the bytes of this page.
 */
public byte[] getPageData() {
    int len = BufferPool.getPageSize();
    ByteArrayOutputStream baos = new ByteArrayOutputStream(len);
    DataOutputStream dos = new DataOutputStream(baos);

    // create the header of the page
    for (int i = 0; i < header.length; i++) {
        try {
            dos.writeByte(header[i]);
        } catch (IOException e) {
            // this really shouldn't happen
            e.printStackTrace();
        }
    }

    // create the tuples
    for (int i = 0; i < tuples.length; i++) {

        // empty slot
        if (!isSlotUsed(i)) {
            for (int j = 0; j < td.getSize(); j++) {
                try {
                    dos.writeByte(0);
                } catch (IOException e) {
                    e.printStackTrace();
                }
            }
        }
    }
}

```

```

        }
        continue;
    }

    // non-empty slot
    for (int j = 0; j < td.numFields(); j++) {
        Field f = tuples[i].getField(j);
        try {
            f.serialize(dos);

        } catch (IOException e) {
            e.printStackTrace();
        }
    }
}

// padding
int zerolen = BufferPool.getPageSize() - (header.length + td.getSize() *
tuples.length); // - numSlots *

        // td.getSize();
        byte[] zeroes = new byte[zerolen];
        try {
            dos.write(zeroes, 0, zerolen);
        } catch (IOException e) {
            e.printStackTrace();
        }

        try {
            dos.flush();
        } catch (IOException e) {
            e.printStackTrace();
        }

        return baos.toByteArray();
    }

/**
 * Static method to generate a byte array corresponding to an empty
 * HeapPage.
 * Used to add new, empty pages to the file. Passing the results of
 * this method to the HeapPage constructor will create a HeapPage with
 * no valid tuples in it.
 *
 * @return The returned ByteArray.

```

```

*/
public static byte[] createEmptyPageData() {
    int len = BufferPool.getPageSize();
    return new byte[len]; // all 0
}

/**
 * Delete the specified tuple from the page; the tuple should be updated to
 * reflect
 * that it is no longer stored on any page.
 *
 * @throws DbException if this tuple is not on this page, or tuple slot is
 * already empty.
 * @param t The tuple to delete
 */
public void deleteTuple(Tuple t) throws DbException {
    // some code goes here
    // not necessary for Lab1
}

/**
 * Adds the specified tuple to the page; the tuple should be updated to reflect
 * that it is now stored on this page.
 *
 * @throws DbException if the page is full (no empty slots) or tupledesc
 * is mismatch.
 * @param t The tuple to add.
 */
public void insertTuple(Tuple t) throws DbException {
    // some code goes here
    // not necessary for Lab1
}

/**
 * Marks this page as dirty/not dirty and record that transaction
 * that did the dirtying
 */
public void markDirty(boolean dirty, TransactionId tid) {
    // some code goes here
    // not necessary for Lab1
}

/**
 * Returns the tid of the transaction that last dirtied this page, or null if
 * the page is not dirty

```

```

*/
public TransactionId isDirty() {
    // some code goes here
    // Not necessary for Lab1
    return null;
}

/**
 * Returns the number of empty slots on this page.
 */
public int getNumEmptySlots() {
    // some code goes here
    int count = 0;
    for (int i = 0; i < tuples.length; i++) {
        if (!isSlotUsed(i)) {
            count++;
        }
    }
    return count;
}

/**
 * Returns true if associated slot on this page is filled.
 */
public boolean isSlotUsed(int i) {
    // some code goes here
    return (header[i / 8] & (1 << (i % 8))) != 0;
}

/**
 * Abstraction to fill or clear a slot on this page.
 */
private void markSlotUsed(int i, boolean value) {
    // some code goes here
    // not necessary for Lab1
}

/**
 * @return an iterator over all tuples on this page (calling remove on this
 *         iterator throws an UnsupportedOperationException)
 *         (note that this iterator shouldn't return tuples in empty slots!)
 */
public Iterator<Tuple> iterator() {
    // some code goes here
    return new HeapPageIterator();
}

```

```

}

private class HeapPageIterator implements Iterator<Tuple> {

    private int curr = 0;

    public boolean hasNext() {
        return curr < tuples.length && isSlotUsed(curr);
    }

    public Tuple next() {
        if (!hasNext()) {
            throw new NoSuchElementException();
        }
        Tuple t = tuples[curr];
        curr++;
        return t;
    }

    public void remove() {
        throw new UnsupportedOperationException();
    }
}
}

```

五、附录

尽管只要求做到 Exercise 4，但实际上全部代码都已经完成，也通过了全部单元测试。这里仅给出这些代码。

(1) Exercise 5

```
// src/java/HeapFile.java
```

```

package simplifiedb;

import java.io.*;
import java.security.Permissions;
import java.util.*;

/**
 * HeapFile is an implementation of a DbFile that stores a collection of tuples
 * in no particular order. Tuples are stored on pages, each of which is a fixed
 * size, and the file is simply a collection of those pages. HeapFile works
 * closely with HeapPage. The format of HeapPages is described in the HeapPage
 * constructor.
 */

```

```

* @see simpledb.HeapPage#HeapPage
* @author Sam Madden
*/
public class HeapFile implements DbFile {

    private final File file;
    private final TupleDesc tuDesc;

    /**
     * Constructs a heap file backed by the specified file.
     *
     * @param f
     *         the file that stores the on-disk backing store for this heap
     *         file.
     */
    public HeapFile(File f, TupleDesc td) {
        // some code goes here
        file = f;
        tuDesc = td;
    }

    /**
     * Returns the File backing this HeapFile on disk.
     *
     * @return the File backing this HeapFile on disk.
     */
    public File getFile() {
        // some code goes here
        return file;
    }

    /**
     * Returns an ID uniquely identifying this HeapFile. Implementation note:
     * you will need to generate this tableid somewhere ensure that each
     * HeapFile has a "unique id," and that you always return the same value for
     * a particular HeapFile. We suggest hashing the absolute file name of the
     * file underlying the heapfile, i.e. f.getAbsolutePath().hashCode().
     *
     * @return an ID uniquely identifying this HeapFile.
     */
    public int getId() {
        // some code goes here
        return file.getAbsolutePath().hashCode();
    }
}

```

```

/**
 * Returns the TupleDesc of the table stored in this DbFile.
 *
 * @return TupleDesc of this DbFile.
 */
public TupleDesc getTupleDesc() {
    // some code goes here
    return tuDesc;
}

// see DbFile.java for javadocs
public Page readPage(PageId pid) {
    // some code goes here
    try {
        RandomAccessFile raf = new RandomAccessFile(file, "r");
        raf.seek(pid.pageNumber() * BufferPool.PAGE_SIZE);
        byte[] data = new byte[BufferPool.PAGE_SIZE];
        raf.read(data);
        raf.close();
        return new HeapPage(new HeapPageId(pid.getTableId(), pid.pageNumber()),
data);
    } catch (IOException e) {
        throw new IllegalArgumentException("Page does not exist");
    }
}

// see DbFile.java for javadocs
public void writePage(Page page) throws IOException {
    // some code goes here
    // not necessary for Lab1
}

/**
 * Returns the number of pages in this HeapFile.
 */
public int numPages() {
    // some code goes here
    return (int) (file.length() / BufferPool.PAGE_SIZE);
}

// see DbFile.java for javadocs
public ArrayList<Page> insertTuple(TransactionId tid, Tuple t)
    throws DbException, IOException, TransactionAbortedException {
    // some code goes here
    return null;
}

```



```

        // not necessary for Lab1
    }

    // see DbFile.java for javadocs
    public ArrayList<Page> deleteTuple(TransactionId tid, Tuple t) throws
    DbException,
        TransactionAbortedException {
        // some code goes here
        return null;
        // not necessary for Lab1
    }

    // see DbFile.java for javadocs
    public DbFileIterator iterator(TransactionId tid) {
        // some code goes here
        return new HeapFileIterator(tid);
    }

    private class HeapFileIterator implements DbFileIterator {

        private final TransactionId tid;
        private int curPage = 0;
        private HeapPage curPageData = null;
        private Iterator<Tuple> curPageIterator = null;

        public HeapFileIterator(TransactionId tid) {
            // some code goes here
            this.tid = tid;
        }

        @Override
        public void open() throws DbException, TransactionAbortedException,
        IOException {
            curPage = 0;
            curPageData = (HeapPage) Database.getBufferPool().getPage(
                tid,
                new HeapPageId(getId(), curPage),
                simpledb.Permissions.READ_WRITE);
            curPageIterator = curPageData.iterator();
        }

        @Override
        public boolean hasNext() throws DbException, TransactionAbortedException,
        IOException {
            if (curPageIterator == null) {

```

```

        return false;
    }
    if (curPageIterator.hasNext()) {
        return true;
    } else {
        if (curPage < numPages() - 1) {
            curPage++;
            curPageData = (HeapPage) Database.getBufferPool().getPage(
                tid,
                new HeapPageId(getId(), curPage),
                simpledb.Permissions.READ_WRITE);
            curPageIterator = curPageData.iterator();
            return hasNext();
        } else {
            return false;
        }
    }
}

@Override
public Tuple next() throws DbException, TransactionAbortedException {
    if (curPageIterator == null) {
        throw new NoSuchElementException("Iterator is not open");
    }
    return curPageIterator.next();
}

@Override
public void rewind() throws DbException, TransactionAbortedException,
IOException {
    curPage = 0;
    curPageData = (HeapPage) Database.getBufferPool().getPage(
        tid,
        new HeapPageId(getId(), curPage),
        simpledb.Permissions.READ_WRITE);
    curPageIterator = curPageData.iterator();
}

@Override
public void close() {
    curPageData = null;
    curPageIterator = null;
}
}
}

```

(2) Exercise 6

```
// src/java/SeqScan.java
```

```
package simplifiedb;
```

```
import java.io.IOException;
```

```
import java.util.*;
```

```
/**
```

```
 * SeqScan is an implementation of a sequential scan access method that reads  
 * each tuple of a table in no particular order (e.g., as they are laid out on  
 * disk).
```

```
 */
```

```
public class SeqScan implements DbIterator {
```

```
    private static final long serialVersionUID = 1L;
```

```
    private final TransactionId tid;
```

```
    private int tableid;
```

```
    private String tableAlias;
```

```
    private DbFileIterator fileIter;
```

```
    /**
```

```
     * Creates a sequential scan over the specified table as a part of the  
     * specified transaction.
```

```
     *
```

```
     * @param tid
```

```
     *           The transaction this scan is running as a part of.
```

```
     * @param tableid
```

```
     *           the table to scan.
```

```
     * @param tableAlias
```

```
     *           the alias of this table (needed by the parser); the  
     *           returned
```

```
     *           tupleDesc should have fields with name tableAlias.fieldName  
     *           (note: this class is not responsible for handling a case  
     *           where
```

```
     *           tableAlias or fieldName are null. It shouldn't crash if  
     *           they
```

```
     *           are, but the resulting name can be null.fieldName,
```

```
     *           tableAlias.null, or null.null).
```

```
    */
```

```
    public SeqScan(TransactionId tid, int tableid, String tableAlias) {
```

```
        // some code goes here
```

```
        this.tid = tid;
```

```

        this.tableid = tableid;
        this.tableAlias = tableAlias;
    }

    /**
     * @return
     *      return the table name of the table the operator scans. This should
     *      be the actual name of the table in the catalog of the database
     */
    public String getTableName() {
        // some code goes here
        return Database.getCatalog().getTableName(tableid);
    }

    /**
     * @return Return the alias of the table this operator scans.
     */
    public String getAlias() {
        // some code goes here
        return tableAlias;
    }

    /**
     * Reset the tableid, and tableAlias of this operator.
     *
     * @param tableid
     *      the table to scan.
     * @param tableAlias
     *      the alias of this table (needed by the parser); the
     *      returned
     *      tupleDesc should have fields with name tableAlias.fieldName
     *      (note: this class is not responsible for handling a case
     *      where
     *      tableAlias or fieldName are null. It shouldn't crash if
     *      they
     *      are, but the resulting name can be null.fieldName,
     *      tableAlias.null, or null.null).
     */
    public void reset(int tableid, String tableAlias) {
        // some code goes here
        this.tableid = tableid;
        this.tableAlias = tableAlias;
    }

    public SeqScan(TransactionId tid, int tableid) {

```

```

        this(tid, tableid, Database.getCatalog().getTableNames(tableid));
    }

    public void open() throws DbException, TransactionAbortedException, IOException
    {
        // some code goes here
        fileIter = Database.getCatalog().getDatabaseFile(tableid).iterator(tid);
        fileIter.open();
    }

    /**
     * Returns the TupleDesc with field names from the underlying HeapFile,
     * prefixed with the tableAlias string from the constructor. This prefix
     * becomes useful when joining tables containing a field(s) with the same
     * name.
     *
     * @return the TupleDesc with field names from the underlying HeapFile,
     *         prefixed with the tableAlias string from the constructor.
     */
    public TupleDesc getTupleDesc() {
        // some code goes here
        TupleDesc tupleDesc = Database.getCatalog().getTupleDesc(tableid);
        Type[] fieldTypes = new Type[tupleDesc.numFields()];
        String[] fieldNames = new String[tupleDesc.numFields()];
        for (int i = 0; i < tupleDesc.numFields(); i++) {
            fieldTypes[i] = tupleDesc.getFieldType(i);
            fieldNames[i] = tableAlias + "." + tupleDesc.getFieldName(i);
        }
        return new TupleDesc(fieldTypes, fieldNames);
    }

    public boolean hasNext() throws TransactionAbortedException, DbException,
    IOException {
        // some code goes here
        return fileIter.hasNext();
    }

    public Tuple next() throws NoSuchElementException,
        TransactionAbortedException, DbException {
        // some code goes here
        return fileIter.next();
    }

    public void close() {
        // some code goes here
    }

```

```
        fileIter.close();
    }

    public void rewind() throws DbException, NoSuchElementException,
        TransactionAbortedException, IOException {
        // some code goes here
        fileIter.rewind();
    }
}
```